

# Introduction to all forms of parallel computing

Prof Alan Edelman

(18.s192/16.s098)

- All material on <https://github.com/mitmath/Parallel-Computing-Spoke>
  - Suggest bookmarking
- Canvas only used for submitting hws, nothing else

# What will you learn in this class?

use Julia to write

- portable parallel programs on GPUs
- use multithreading on your own laptops
- and distributed computing on multiprocessors.

Get a good feel for what is and is not possible in gaining speedups and performance,

Have a sense of possible research directions including applications and the use of AI, LLMs, and ML, and perhaps most valuably, real progress towards making parallel computing easier.

# What should you learn in another class?


6.1060 performance engineering - low-level optimization

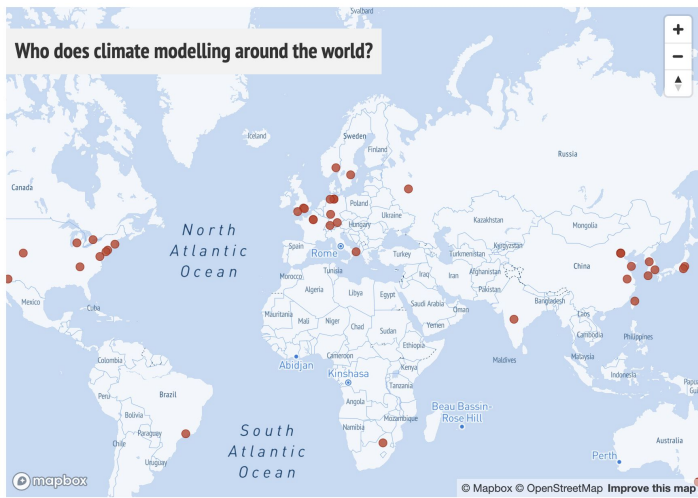
6.S894 low-level performance engineering for GPUs

6.1100 computer language engineering - compiler tinkering


# Why High-Performance Computing and Parallelism

1. You cannot escape: HPC is will be a >100B market in less than a decade





Climate sensitivity in CMIP6 models



2. Large data make faster larger computers necessary

Large Climate models run on supercomputers at big centers

Source:

<https://www.carbonbrief.org/cmip6-the-next-generation-of-climate-models-explained/>

### 3. Because it's fun!



## Parallelism



# Biggest challenge to HPC (my opinion)

- Software situation is really not great (understatement?)
- Too many kinds of parallelism, not enough coherence (yet?)
- Research community tends to prefer performance boasting to usable software
- Same problem gets solved over and over again

# Intro to concepts of parallel computing


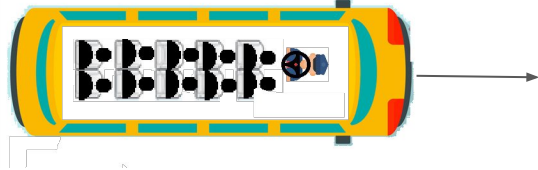
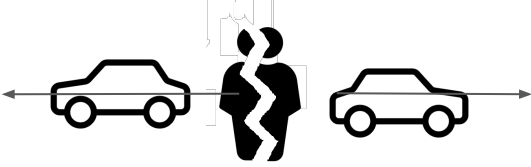
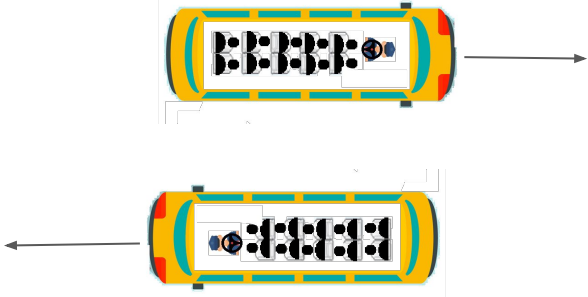


# Flops

Used as #floating point ops and also floating ops/second


- Discussion of top500.org including exaflop machines
- Comparing the rate of matrix multiply computation vs the rate of matrix addition computation.
  - Note: rate means we are counting the speed of each operation,normalizing by the number of operations.

# The concept of parallelism

|                      | Single data   | Multiple data  |
|----------------------|---|--|
| Single instruction   |  |   |
| Multiple instruction |  |  |


# The concept of parallelism

Shared memory



Two bus services serve all stops

Distributed memory




Each bus service has their own hubs; and there is one slower connection between them

# Levels of parallelism

Single core




Multithreading



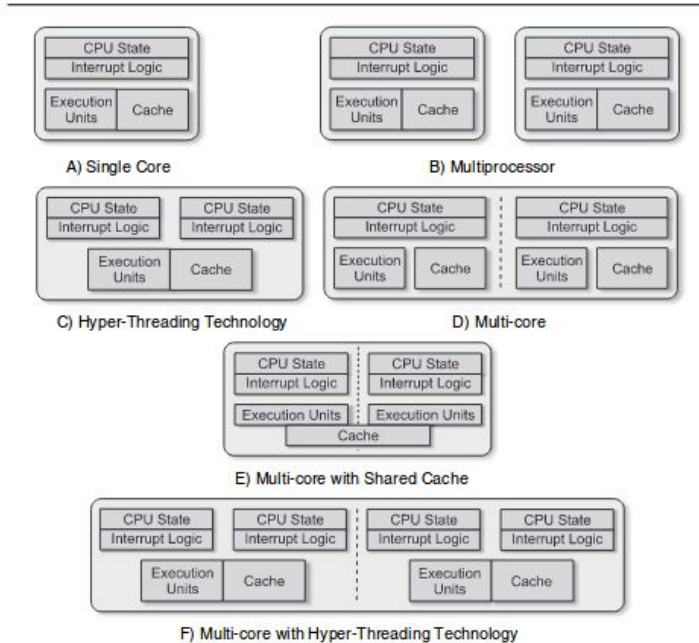
Multicore



Multiprocessor



# Levels of parallelism - the technical and all the options



In julia, multithreading is achieved with tasks

**Figure 1.4** Simple Comparison of Single-core, Multi-processor, and Multi-Core Architectures